

# CSE 502:

# Computer Architecture

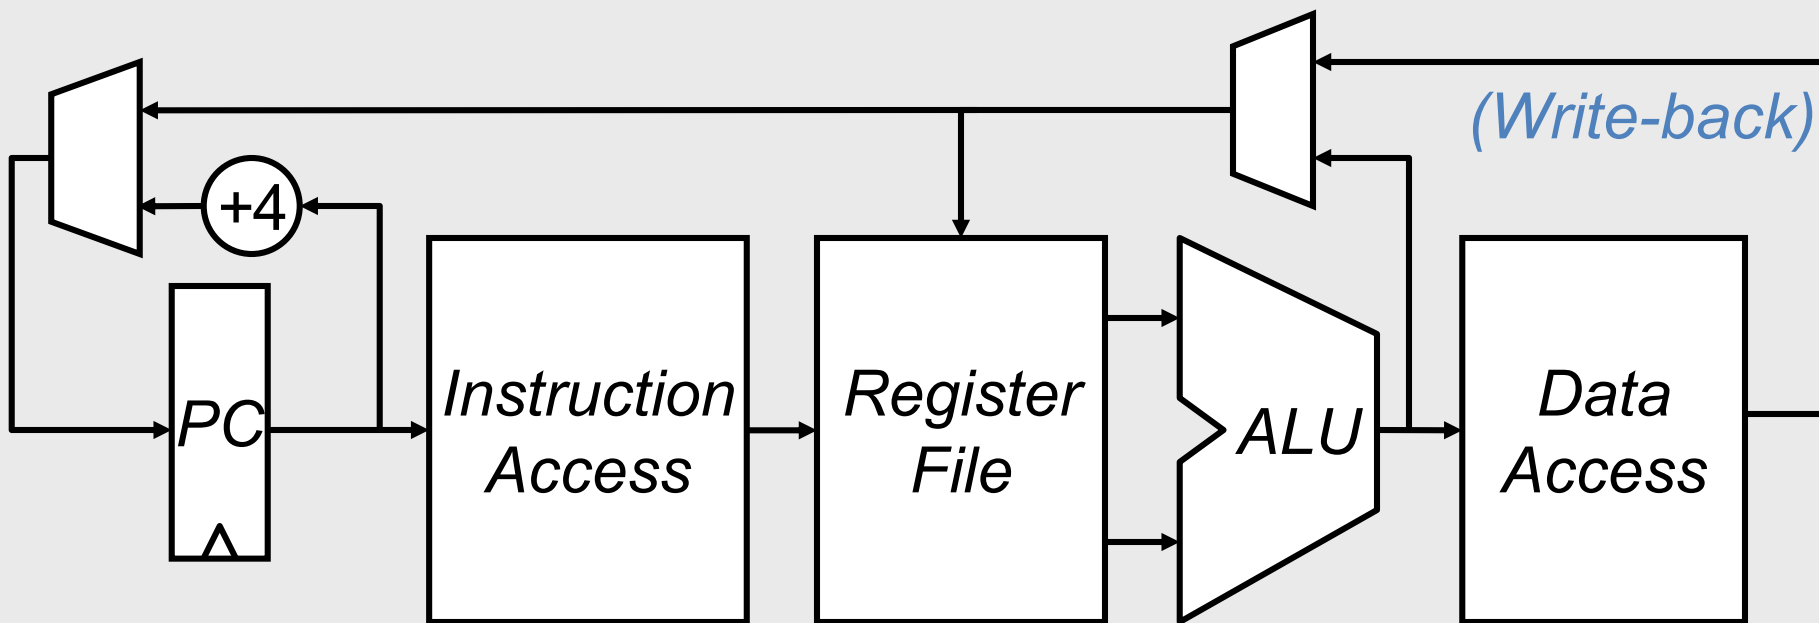
Core Pipelining

# Generic Instruction Cycle

- Steps in processing an instruction:
  - Instruction Fetch (**IF\_STEP**)
  - Instruction Decode (**ID\_STEP**)
  - Operand Fetch (**OF\_STEP**)
  - Execute (**EX\_STEP**)
  - Result Store or Write Back (**RS\_STEP**)
- Actions per instruction at each stage given by ISA
- $\mu$ Arch determines how HW implements the steps

# Prototypical Processor Organization

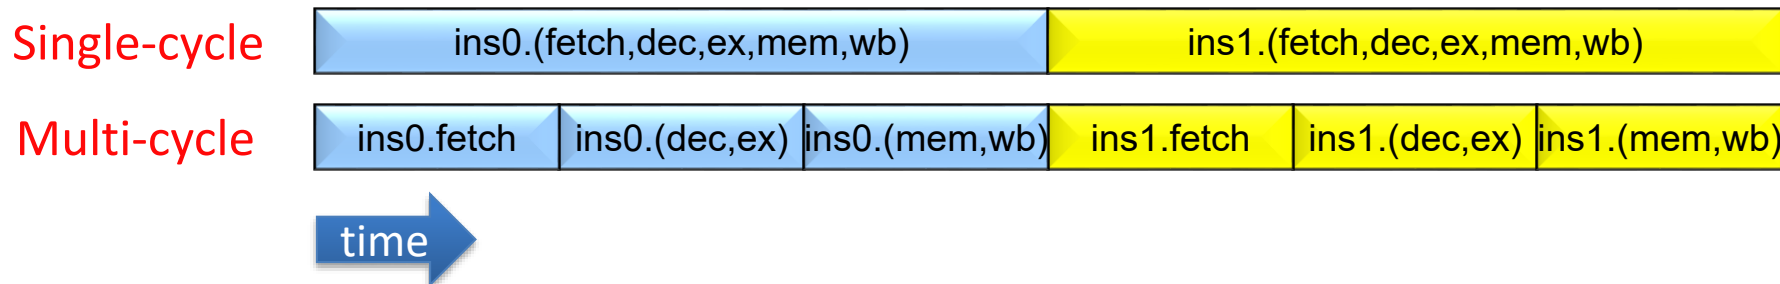
*Addr-gen.    Fetch    Decode    Issue    Execute    Memory*



# Datapath vs. Control Logic

- ***Datapath*** is HW components and connections
  - Determines the *static* structure of processor
- ***Control logic*** controls data flow in datapath
  - Control is determined by
    - Instruction words
    - State of the processor
    - Execution results at each stage

# Single-Instruction Datapath



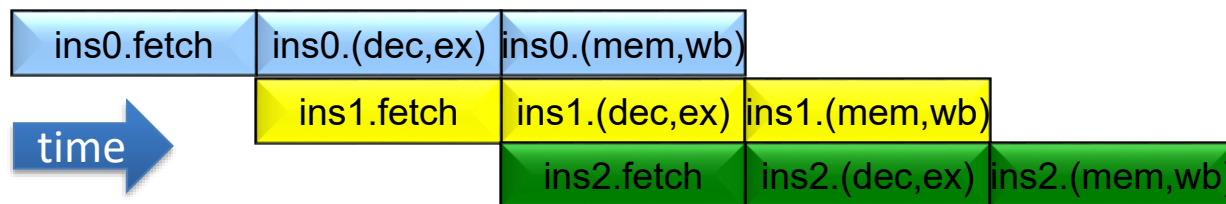
- Process one instruction at a time
- Single-cycle control: hardwired
  - Low CPI (1)
  - Long clock period (to accommodate slowest instruction)
- Multi-cycle control: typically micro-programmed
  - Short clock period
  - High CPI
- Can we have both low CPI and short clock period?
  - Not if datapath executes only one instruction at a time
  - No good way to make a single instruction go faster

# Pipelined Datapath

Multi-cycle



Pipelined




- Start with multi-cycle design
- When insn0 goes from stage 1 to stage 2  
... insn1 starts stage 1
- Each instruction passes through all stages  
... but instructions enter and leave at faster *rate*


Style	Ideal CPI	Cycle Time (1/freq)
Single-cycle	1	Long
Multi-cycle	> 1	Short
<b>Pipelined</b>	<b>1</b>	<b>Short</b>

Pipeline can have as many insns *in flight* as there are stages


# Pipeline Examples



Stage delay =  $n$   
Bandwidth =  $\sim(1/n)$



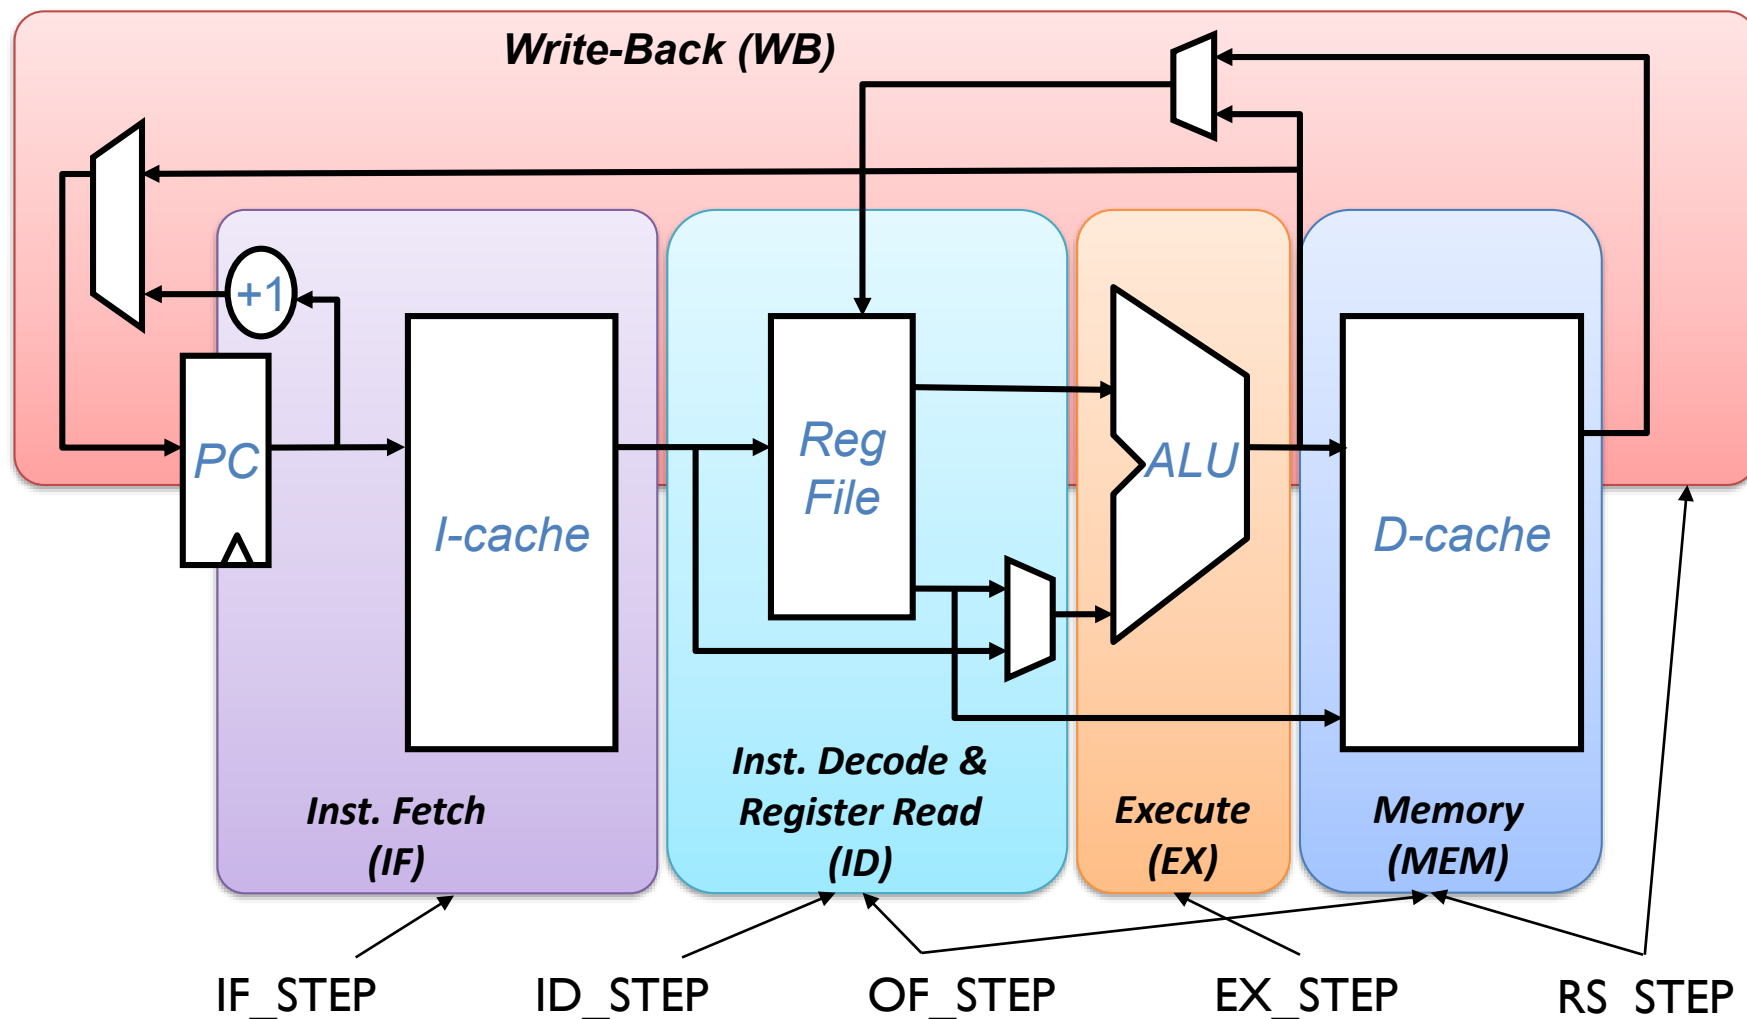
Stage delay =  $n/2$   
Bandwidth =  $\sim(2/n)$



Stage delay =  $n/3$   
Bandwidth =  $\sim(3/n)$

Increases throughput at the expense of latency

# 5-Stage MIPS Datapath

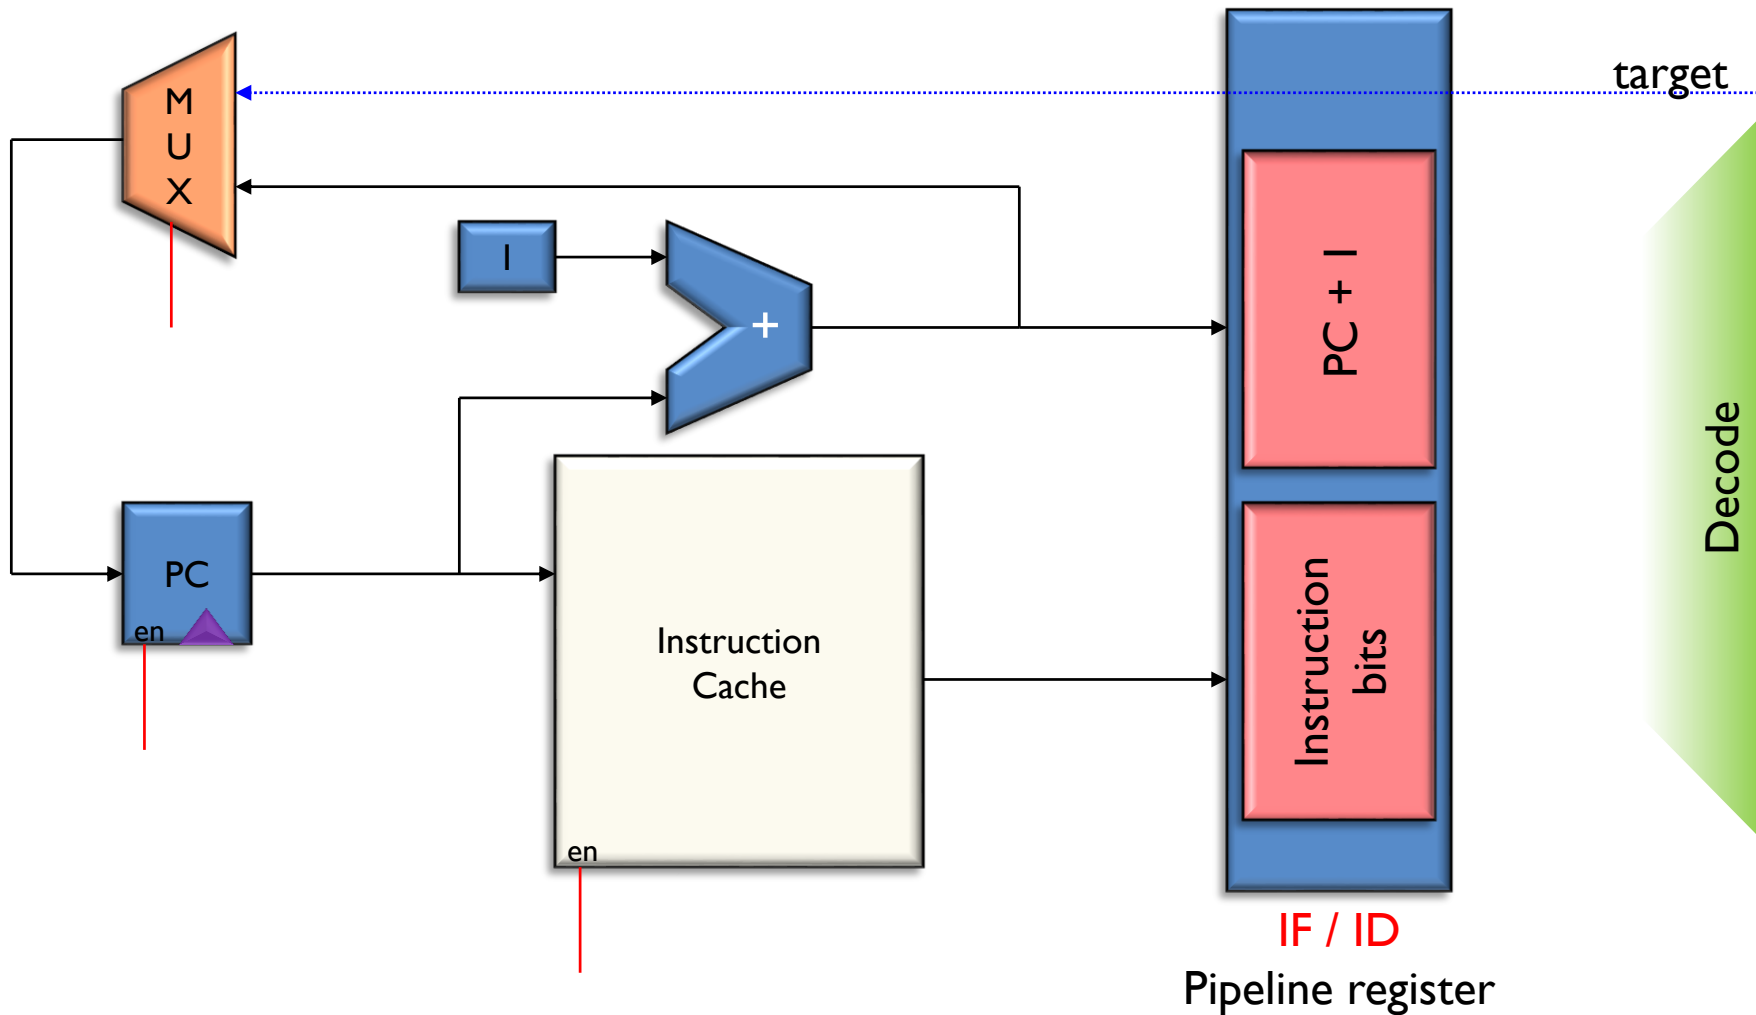




# Stage 1: Fetch

- Fetch instruction from instruction cache
  - Use PC to index instruction cache
  - Increment PC (assume no branches for now)
- Write state to the pipeline register (IF/ID)
  - The next stage will read this pipeline register

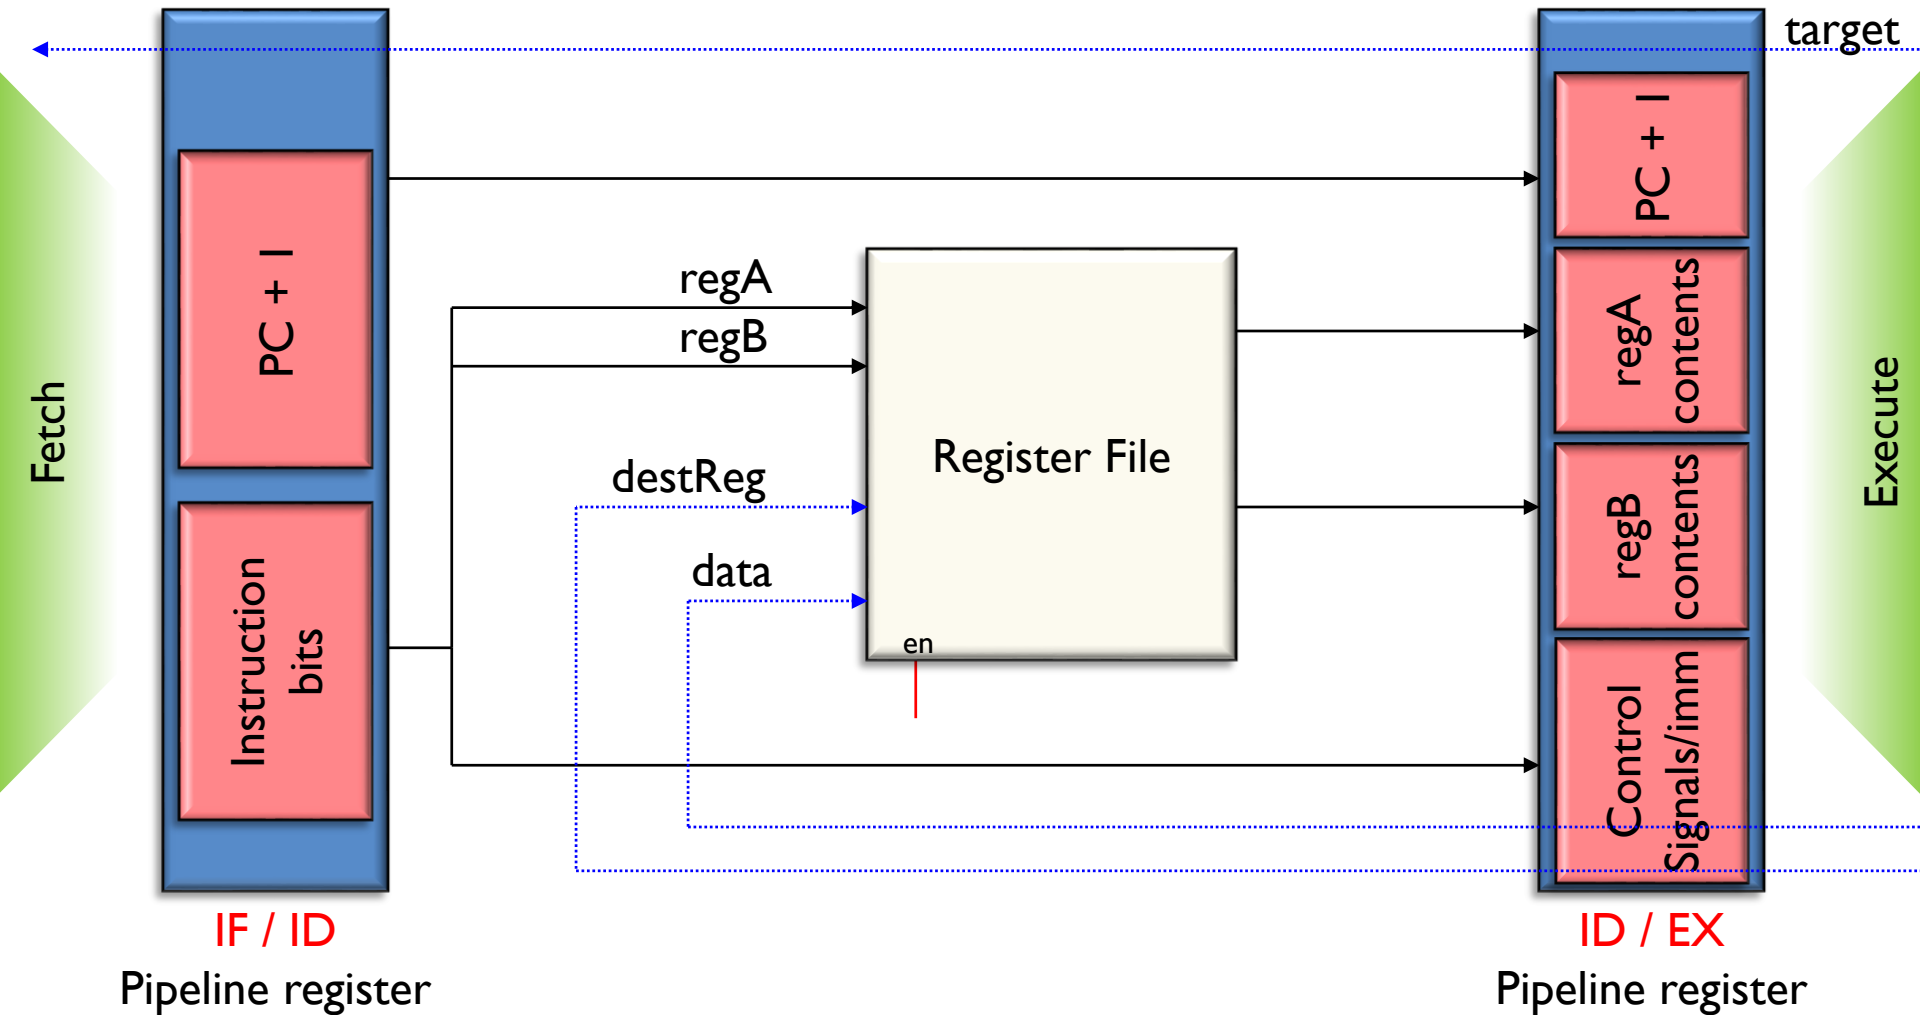
# Stage 1: Fetch Diagram



## Stage 2: Decode

- Decodes opcode bits
  - Set up Control signals for later stages
- Read input operands from register file
  - Specified by decoded instruction bits
- Write state to the pipeline register (ID/EX)
  - Opcode
  - Register contents, immediate operand
  - PC+1 (even though decode didn't use it)
  - Control signals (from insn) for opcode and destReg

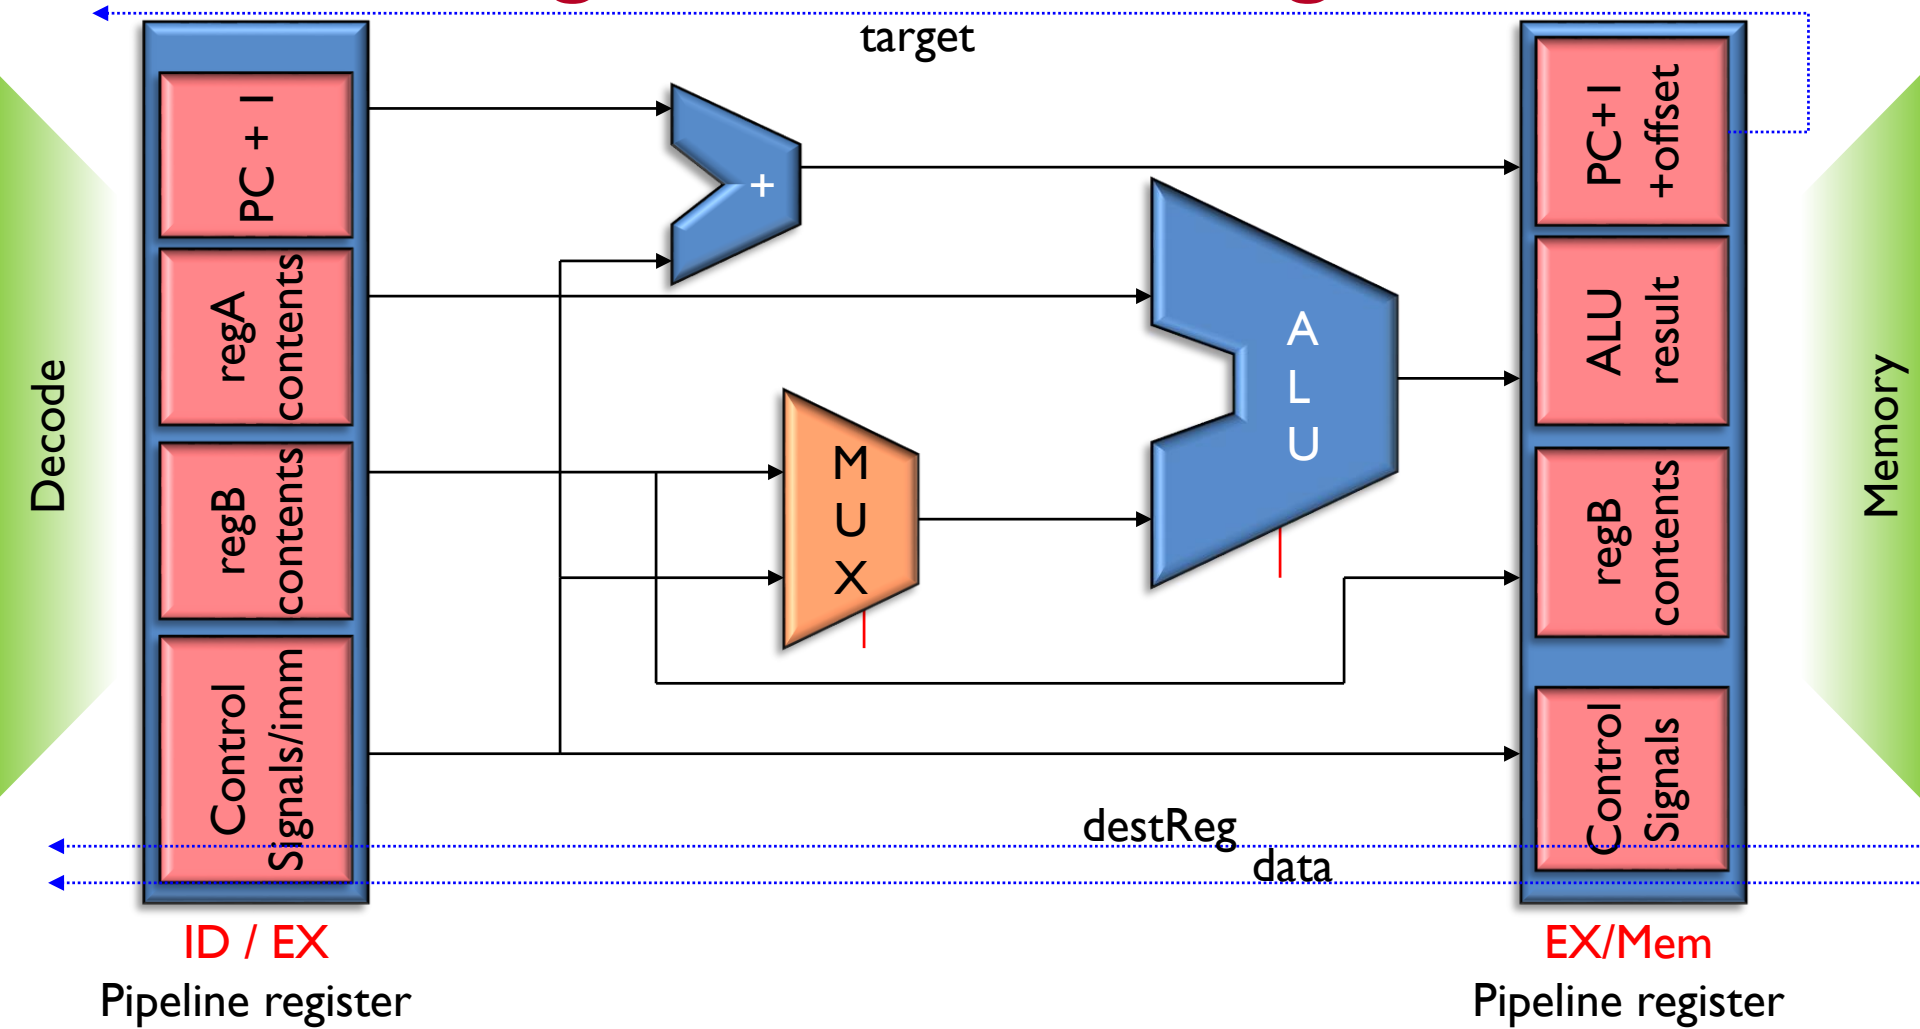
# Stage 2: Decode Diagram



## Stage 3: Execute

- Perform ALU operations
  - Calculate result of instruction
    - Control signals select operation
    - Contents of regA used as one input
    - Either regB or constant offset (imm from insn) used as second input
  - Calculate PC-relative branch target
    - $PC+1+(\text{constant offset})$
- Write state to the pipeline register (EX/Mem)
  - ALU result, contents of regB, and  $PC+1+\text{offset}$
  - Control signals (from insn) for opcode and destReg

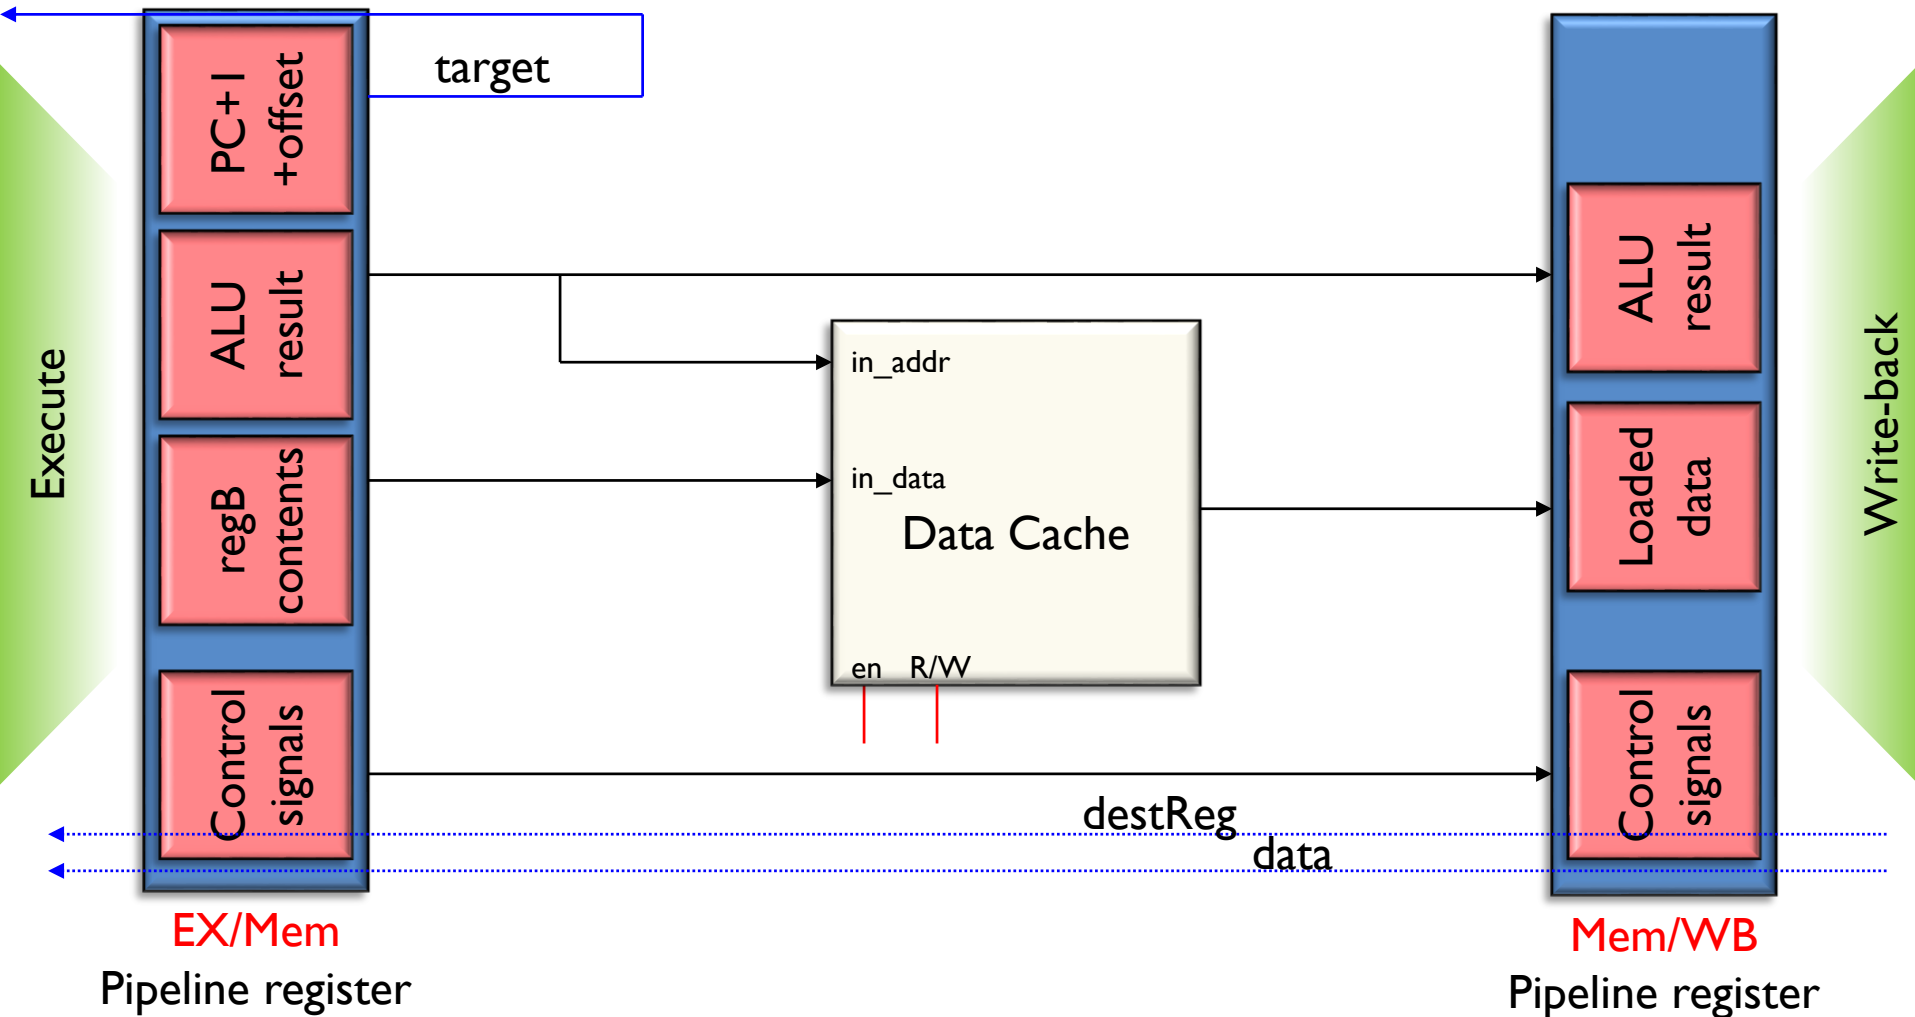
# Stage 3: Execute Diagram



## Stage 4: Memory

- Perform data cache access
  - ALU result contains address for LD or ST
  - Opcode bits control R/W and enable signals
- Write state to the pipeline register (Mem/WB)
  - ALU result and Loaded data
  - Control signals (from insn) for opcode and destReg

# Stage 4: Memory Diagram

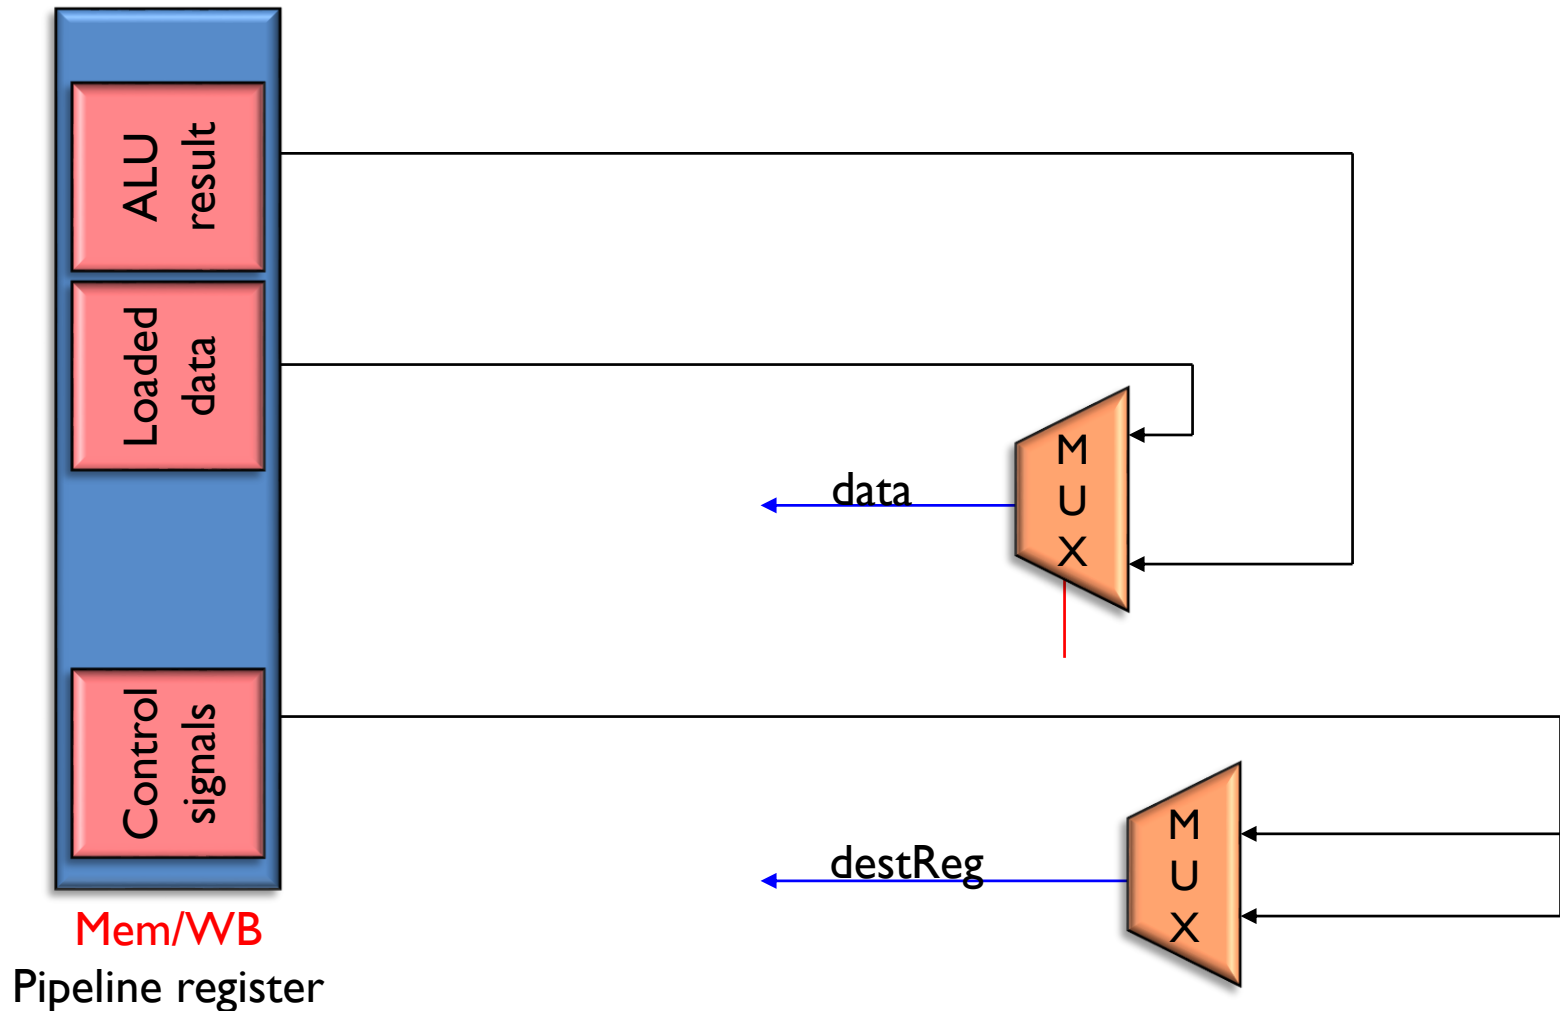




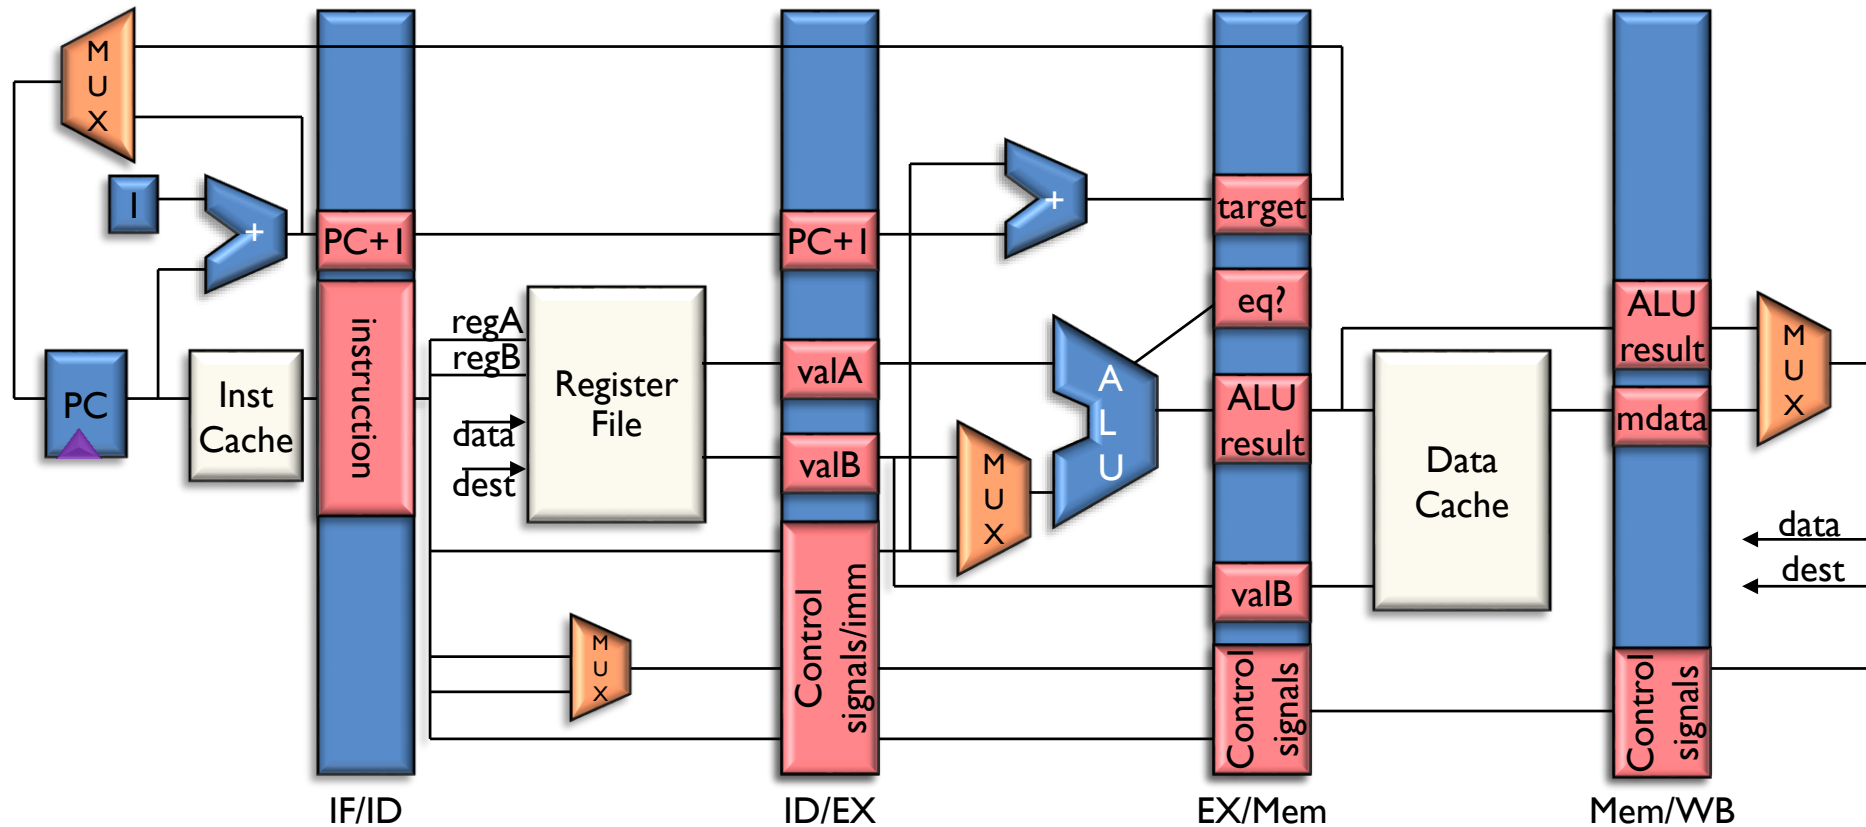
## Stage 5: Write-back

- Writing result to register file (if required)
  - Write Loaded data to destReg for LD
  - Write ALU result to destReg for ALU insn
  - Opcode bits control register write enable signal

# Stage 5: Write-back Diagram



# Putting It All Together



# Pipelining Idealism

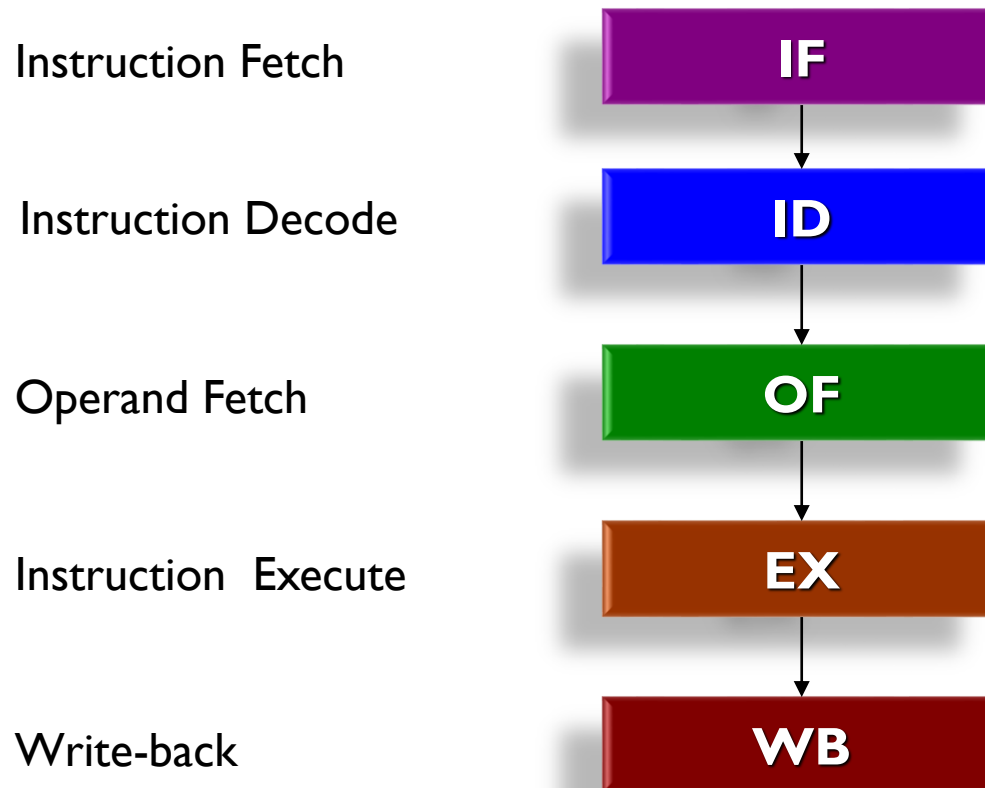
- Uniform Sub-operations
  - Operation can partitioned into uniform-latency sub-ops
- Repetition of Identical Operations
  - Same ops performed on many different inputs
- Independent Operations
  - All ops are mutually independent

# Pipeline Realism

- Uniform Sub-operations ... **NOT!**
  - Balance pipeline stages
    - Stage quantization to yield balanced stages
    - Minimize internal fragmentation (left-over time near end of cycle)
- Repetition of Identical Operations ... **NOT!**
  - Unifying instruction types
    - Coalescing instruction types into one “multi-function” pipe
    - Minimize external fragmentation (idle stages to match length)
- Independent Operations ... **NOT!**
  - Resolve data and resource hazards
    - Inter-instruction dependency detection and resolution

Pipelining is expensive


# The Generic Instruction Pipeline



# Balancing Pipeline Stages



$T_{IF} = 6$  units



$T_{ID} = 2$  units



$T_{OF} = 9$  units



$T_{EX} = 5$  units



$T_{OS} = 9$  units

Without pipelining

$$T_{cyc} \approx T_{IF} + T_{ID} + T_{OF} + T_{EX} + T_{OS} \\ = 31$$

Pipelined

$$T_{cyc} \approx \max\{T_{IF}, T_{ID}, T_{OF}, T_{EX}, T_{OS}\} \\ = 9$$

$$\text{Speedup} = 31 / 9 = 3.44$$

Can we do better?

# Balancing Pipeline Stages (1/2)


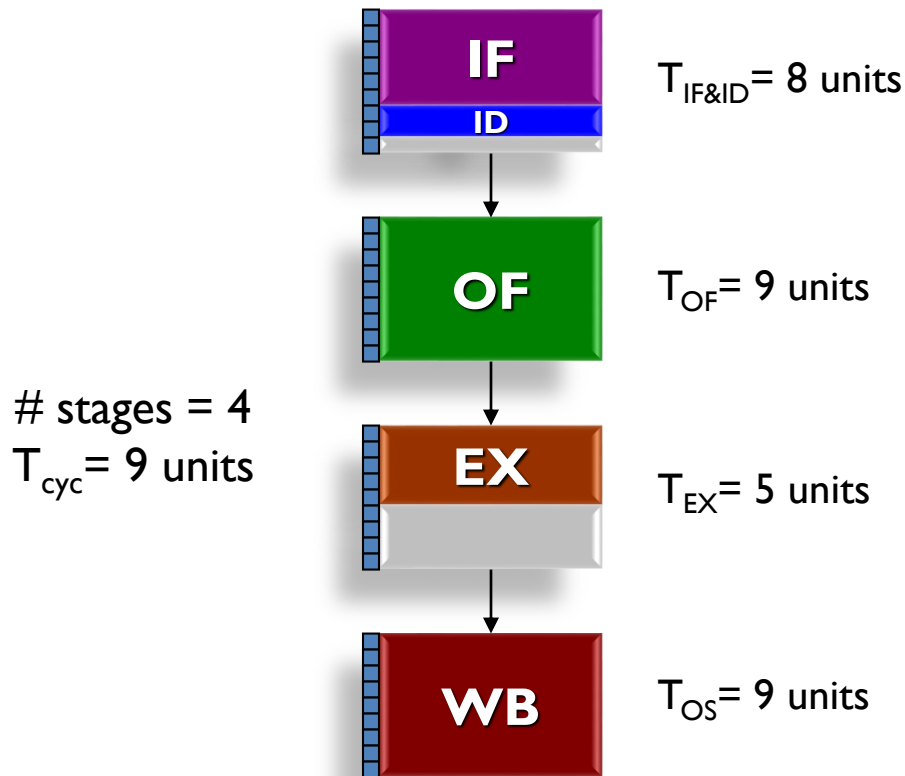
- Two methods for stage quantization
  - Divide sub-ops into smaller pieces
  - Merge multiple sub-ops into one
- Recent/Current trends
  - Deeper pipelines (more and more stages)
  - Pipelining of memory accesses
  - Multiple different pipelines/sub-pipelines



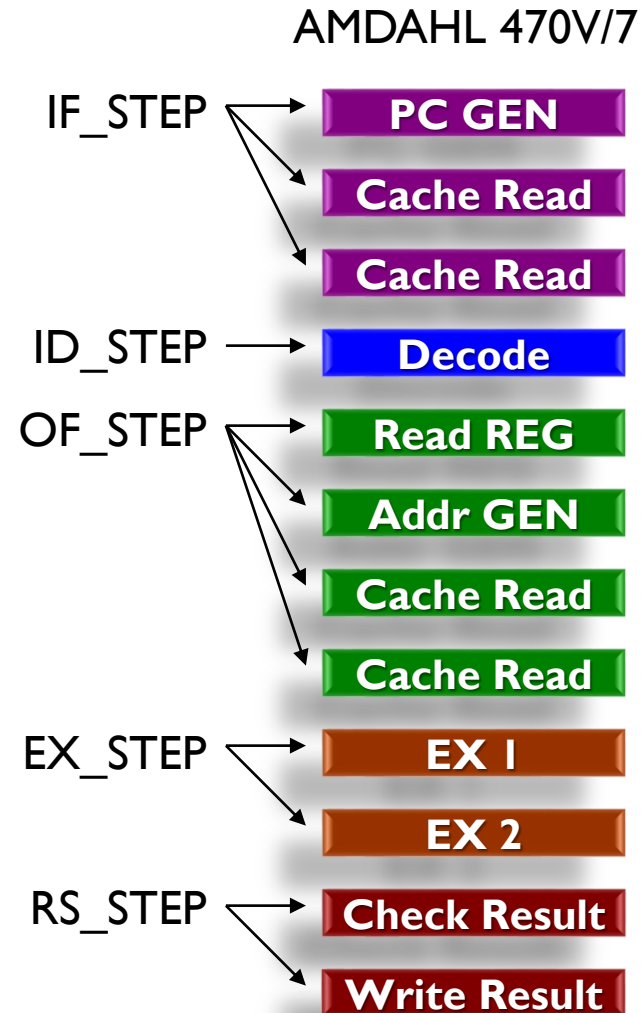
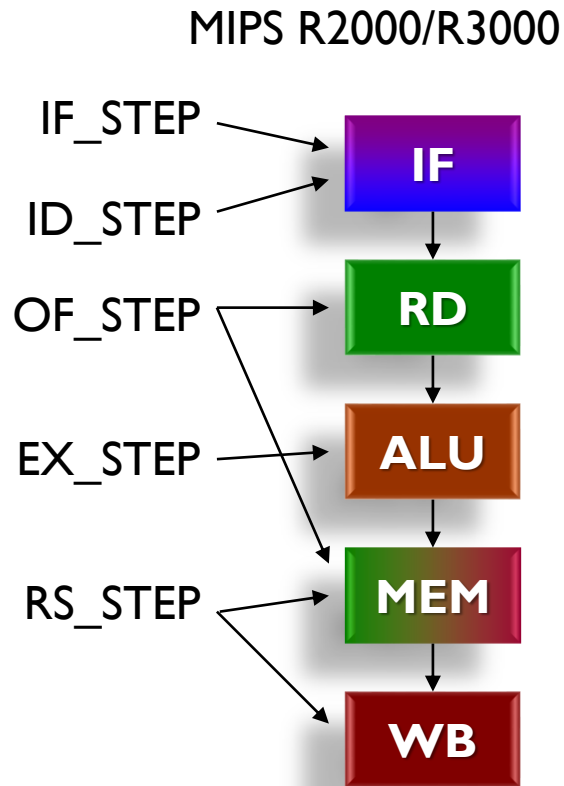
# Balancing Pipeline Stages (2/2)

Coarser-Grained Machine Cycle:  
4 machine cyc / instruction

Finer-Grained Machine Cycle:  
11 machine cyc / instruction



# Pipeline Examples




# Instruction Dependencies (1/2)

- Data Dependence
  - Read-After-Write (RAW) (the only true dependence)
    - Read must wait until earlier write finishes
  - Anti-Dependence (WAR)
    - Write must wait until earlier read finishes (avoid clobbering)
  - Output Dependence (WAW)
    - Earlier write can't overwrite later write
- Control Dependence (a.k.a. Procedural Dependence)
  - Branch condition must execute before branch target
  - Instructions after branch cannot run before branch

# Instruction Dependencies (1/2)

From Quicksort:     # for ( ; (j < high) && (array[j] < array[low]); ++j);



Real code has lots of dependencies


# Hardware Dependency Analysis

- Processor must handle
  - Register Data Dependencies (same register)
    - RAW, WAW, WAR
  - Memory Data Dependencies (same address)
    - RAW, WAW, WAR
  - Control Dependencies

# Pipeline Terminology

- Pipeline Hazards
  - Potential violations of program dependencies
    - Due to multiple in-flight instructions
  - Must ensure program dependencies are not violated
- Hazard Resolution
  - Static method: compiler guarantees correctness
    - By inserting No-Ops or independent insns between dependent insns
  - Dynamic method: hardware checks at runtime
    - Two basic techniques: **Stall** (costs perf.), **Forward** (costs hw)
- Pipeline Interlock
  - Hardware mechanism for dynamic hazard resolution
  - Must detect and enforce dependencies at runtime

# Pipeline: Steady State




# Data Hazards

- Necessary conditions:
  - WAR: write stage earlier than read stage
    - Is this possible in IF-ID-RD-EX-MEM-WB?
  - WAW: write stage earlier than write stage
    - Is this possible in IF-ID-RD-EX-MEM-WB?
  - RAW: read stage earlier than write stage
    - Is this possible in IF-ID-RD-EX-MEM-WB?
- If conditions not met, no need to resolve
- Check for both register and memory

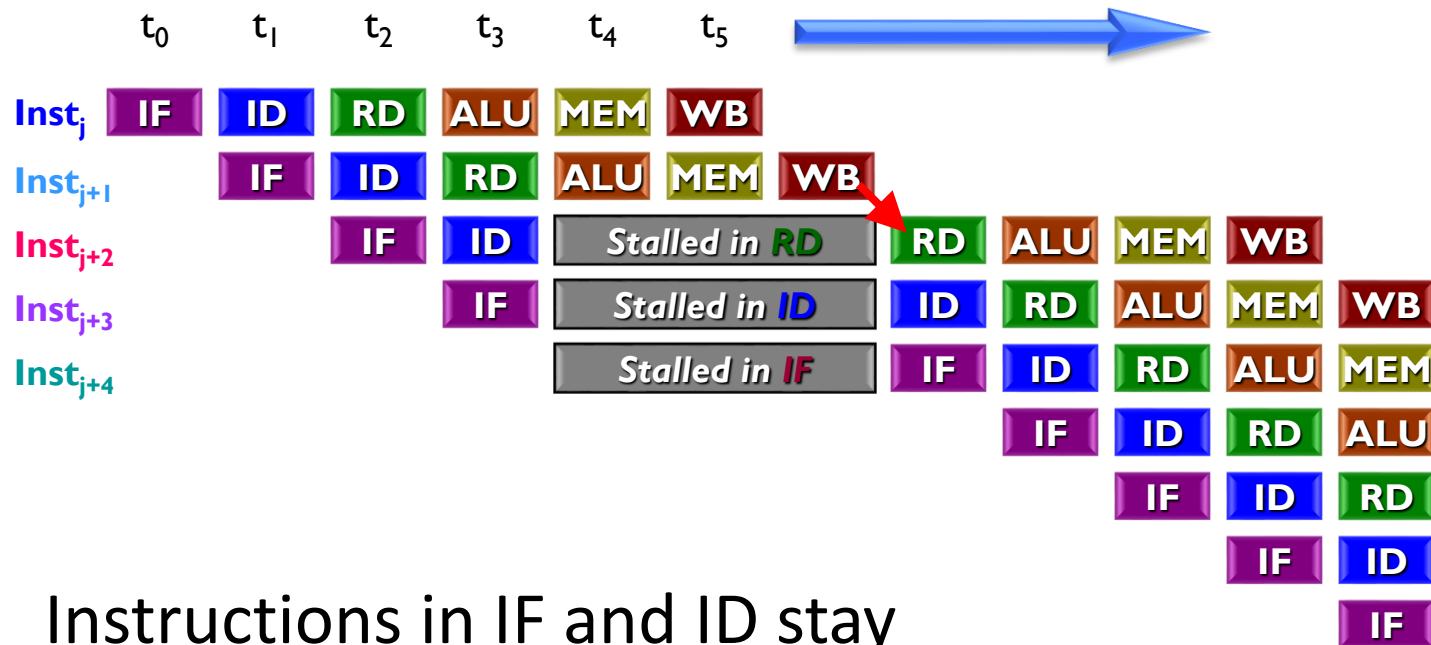


# Pipeline: Data Hazard




- Only RAW in our case
- How to detect?
  - Compare read register specifiers for newer instructions with write register specifiers for older instructions

# Option 1: Stall on Data Hazard



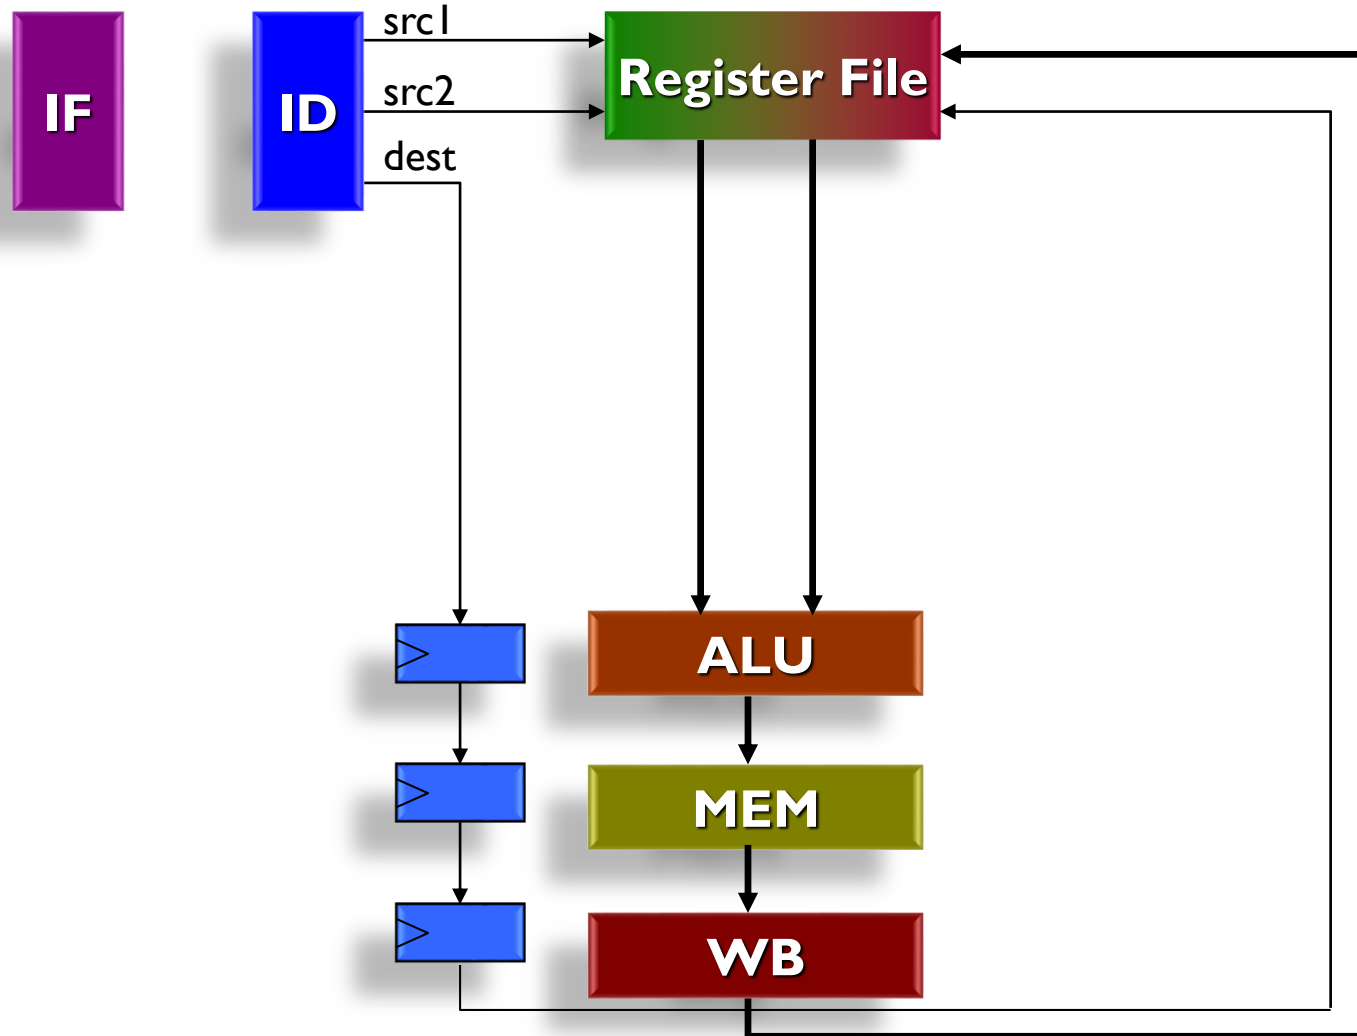
- Instructions in IF and ID stay
- IF/ID pipeline latch not updated
- Send no-op down pipeline (called a bubble)

# Option 2: Forwarding Paths (1/3)

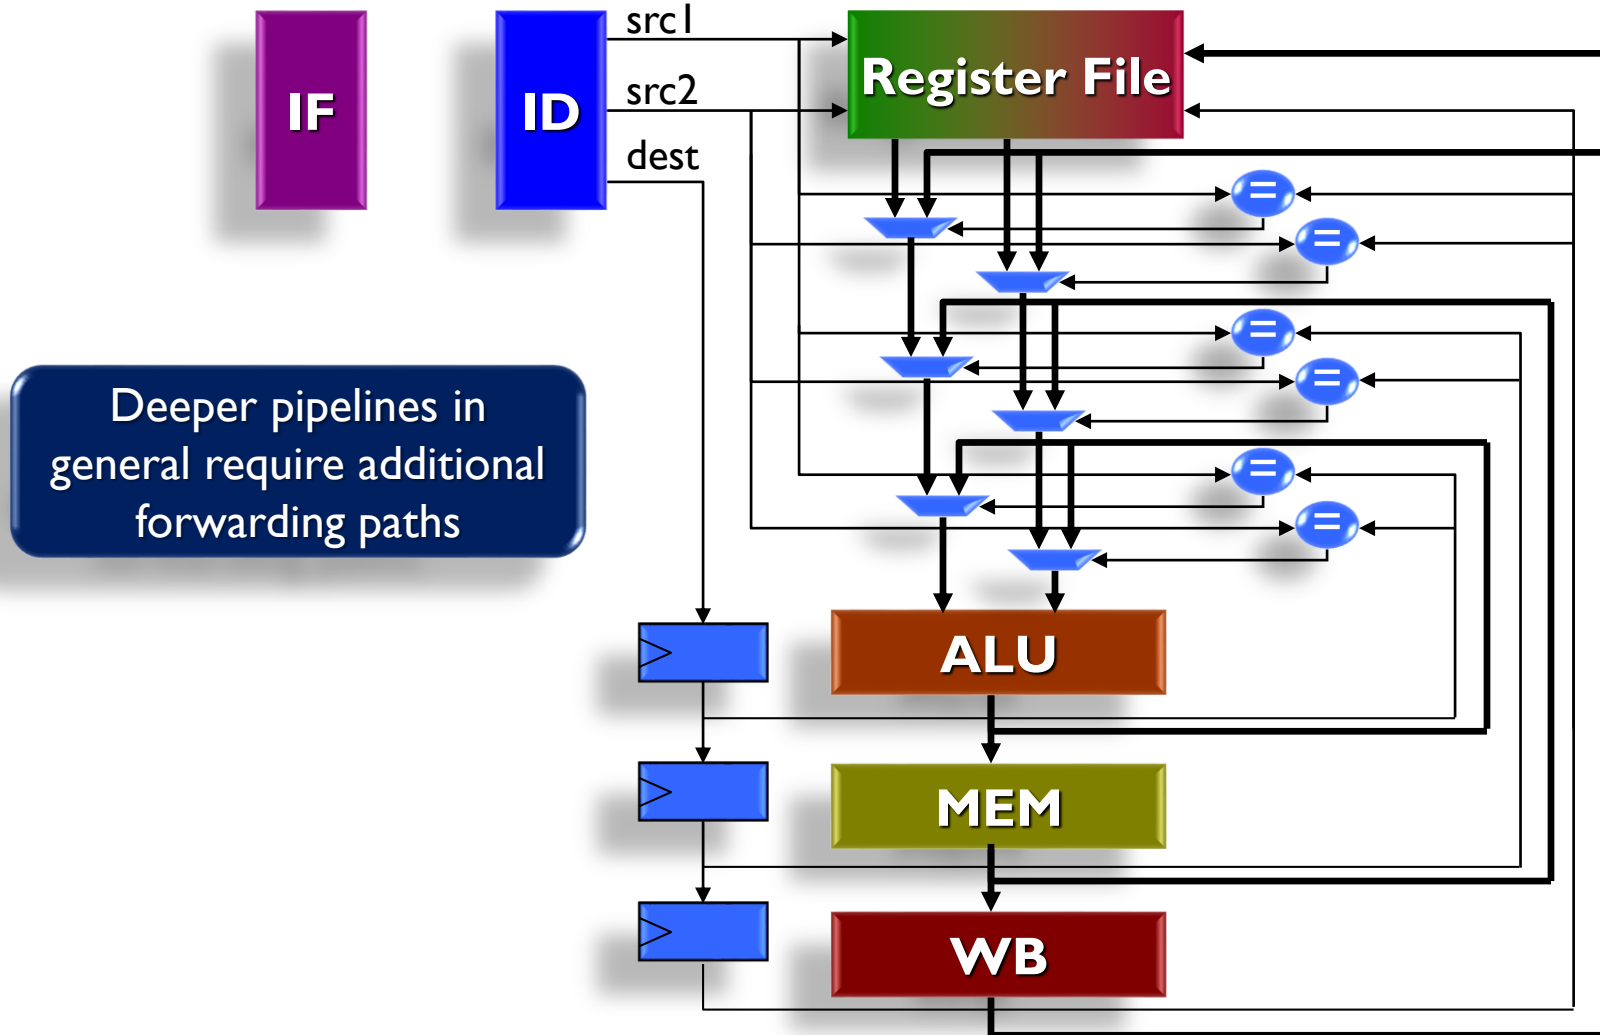


Requires stalling even with forwarding paths


## Option 2: Forwarding Paths (2/3)



# Option 2: Forwarding Paths (3/3)

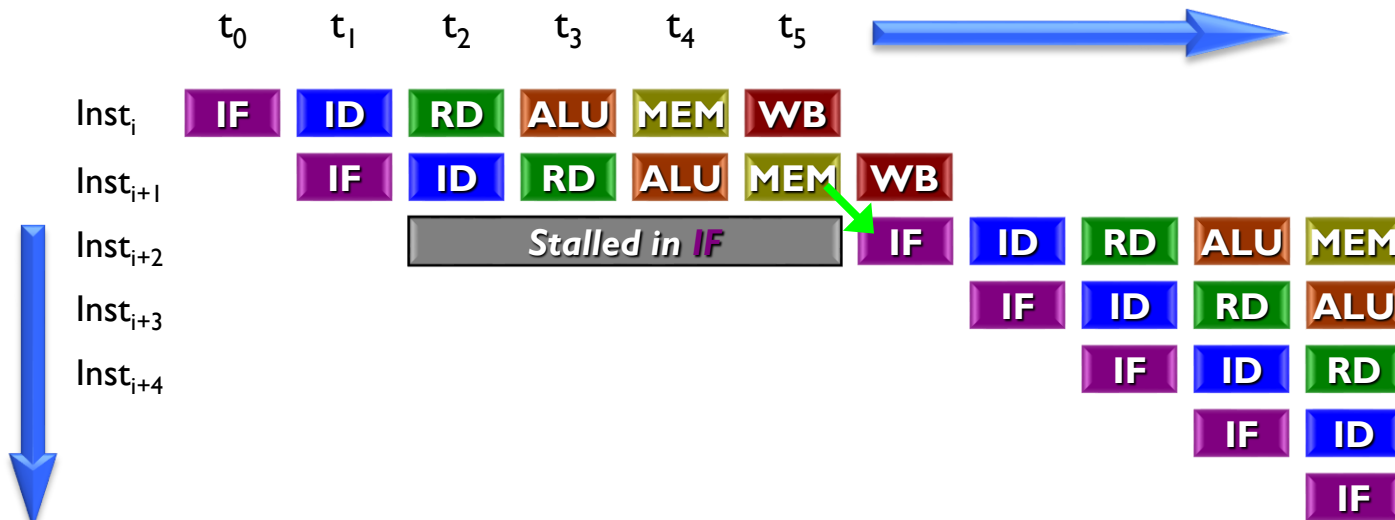


# Pipeline: Control Hazard




- **Note:** The target of  $Inst_{i+1}$  is available at the end of the ALU stage, but it takes one more cycle (MEM) to be written to the PC register

# Option 1: Stall on Control Hazard



- Stop fetching until branch outcome is known
  - Send no-ops down the pipe
- Easy to implement
- Performs poorly
  - ~1 of 6 instructions are branches
  - Each branch takes 4 cycles
  - $CPI = 1 + 4 \times 1/6 = 1.67$  (lower bound)

## Option 2: Prediction for Control Hazards



- Predict branch not taken
- Send sequential instructions down pipeline
- Must stop memory and RF writes
- Kill instructions later if incorrect; we would know at the end of ALU
- Fetch from branch target



## Option 3: Delay Slots for Control Hazards

- Another option: delayed branches
  - # of delay slots ( $ds$ ) : stages between IF and where the branch is resolved
    - 3 in our example
  - Always execute following  $ds$  instructions
  - Put useful instruction there, otherwise no-op
- Losing popularity
  - Just a stopgap (one cycle, one instruction)
  - Superscalar processors (later)
    - Delay slot just gets in the way (special case)


Legacy from old RISC ISAs

# Going Beyond Scalar

- Scalar pipeline limited to  $\text{CPI} \geq 1.0$ 
  - Can never run more than 1 insn per cycle
- “Superscalar” can achieve  $\text{CPI} \leq 1.0$  (i.e.,  $\text{IPC} \geq 1.0$ )
  - Superscalar means executing multiple insns in parallel


# Architectures for Instruction Parallelism

- Scalar pipeline (baseline)
  - Instruction overlap parallelism =  $D$
  - Operation Latency = 1
  - Peak IPC = 1.0

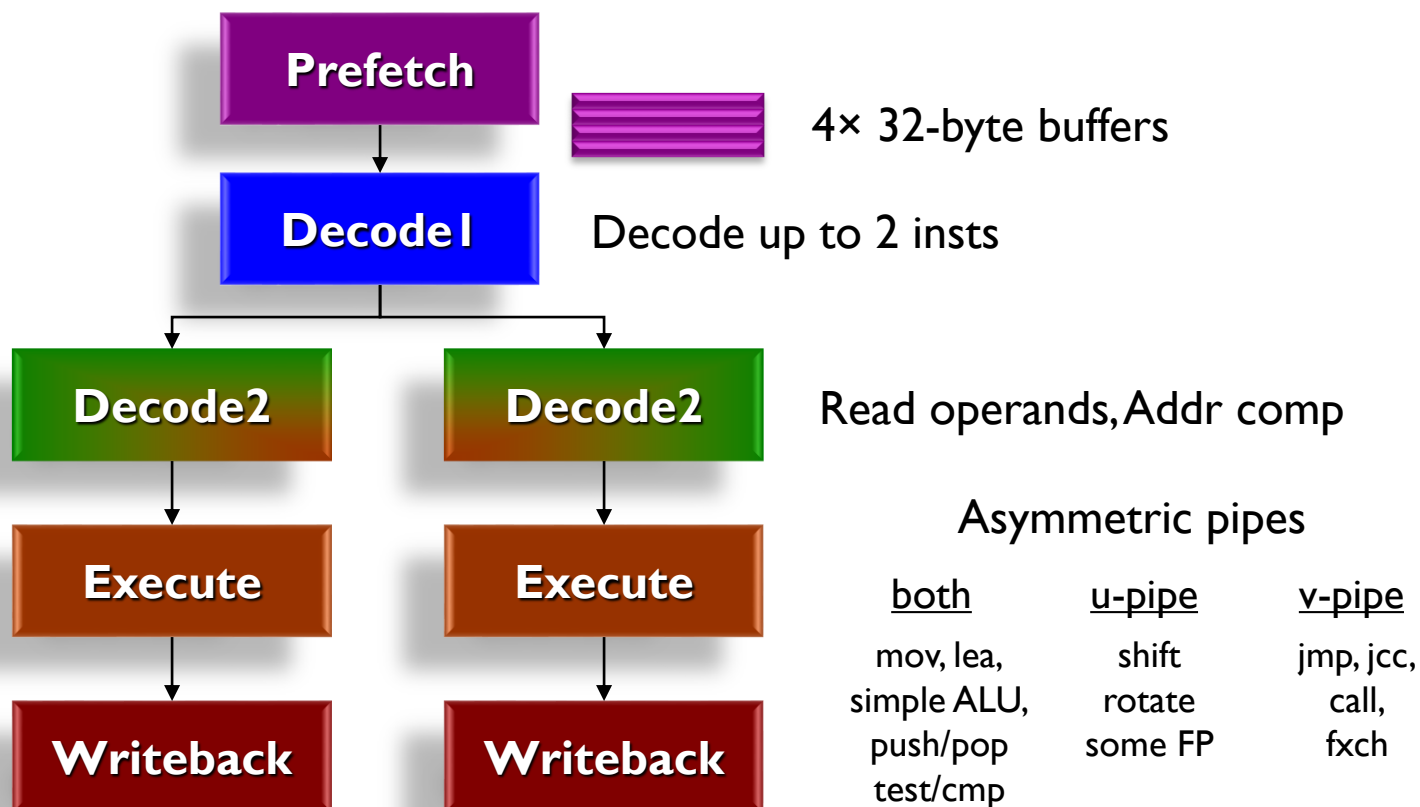


# Superscalar Machine

- Superscalar (pipelined) Execution
  - Instruction parallelism =  $D \times N$
  - Operation Latency = 1
  - Peak IPC =  $N$  per cycle



# Superscalar Example: Pentium



# Pentium Hazards & Stalls

- “Pairing Rules” (when can’t two insns exec?)
  - Read/flow dependence
    - `mov eax, 8`
    - `mov [ebp], eax`
  - Output dependence
    - `mov eax, 8`
    - `mov eax, [ebp]`
  - Partial register stalls
    - `mov al, 1`
    - `mov ah, 0`
  - Function unit rules
    - Some instructions can never be paired
      - `MUL`, `DIV`, `PUSHA`, `MOVS`, some FP

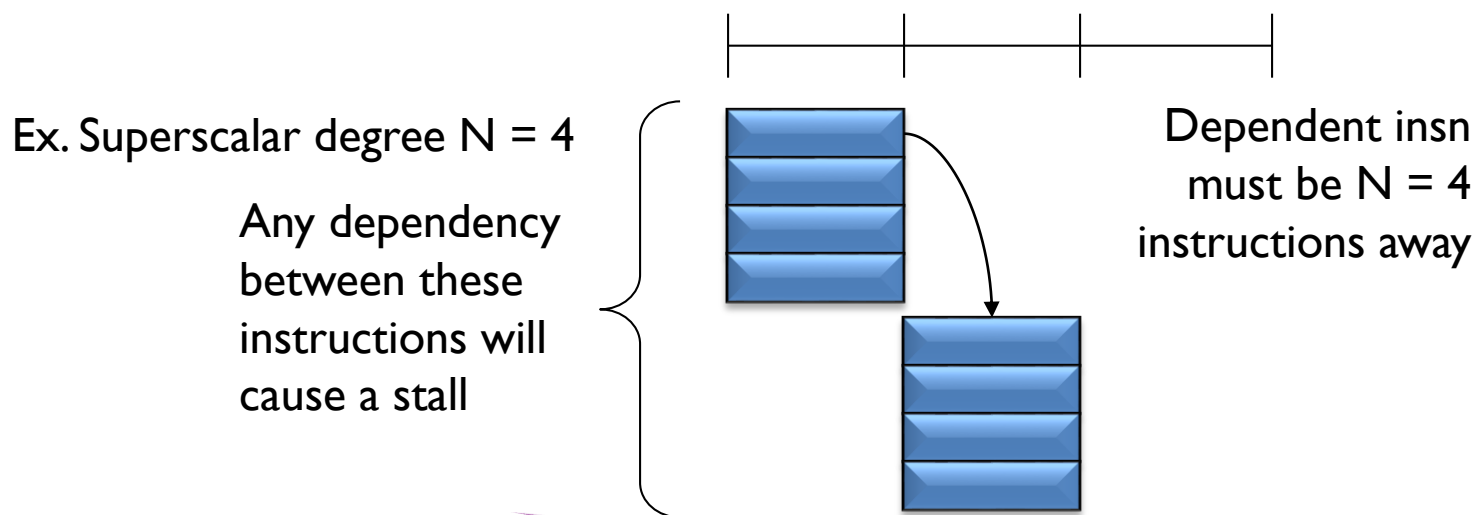
# Limitations of In-Order Pipelines

- If the machine parallelism is increased
  - ... dependencies reduce performance
  - CPI of in-order pipelines degrades sharply
    - As  $N$  approaches avg. distance between dependent instructions
    - Forwarding is no longer effective
  - Must stall often

In-order pipelines are rarely full

# The In-Order N-Instruction Limit

- On average, parent-child separation is about 5 insn
  - (Franklin and Sohi '92)



Average of 5 means there are many cases when the separation is  $< 4$ ...  
each of these limits parallelism

Reasonable in-order superscalar is effectively  $N=2$